

Econometrics I

Lecture 5b: Causal Diagrams (DAGs)

Chris Conlon


NYU Stern

Fall 2026

Two languages for the same problem

- ▶ **Potential outcomes** (last hour) tell us *what we want*: $\mathbb{E}[Y_i(1) - Y_i(0)]$, and what randomization buys us.
- ▶ **Causal diagrams** (DAGs) encode *what we believe*: which variables cause which, drawn as arrows. The assumptions are the *missing* arrows.
- ▶ Why learn both?
 - DAGs make “what should I control for?” a mechanical question with a right answer — and expose when a control makes things **worse**.
 - Potential outcomes handle heterogeneity and identification-by-design more gracefully. Economists mostly publish in this language.
- ▶ Readings: Mixtape Ch. 3; Hünermund & Bareinboim (2021); Cinelli, Forney & Pearl (2022), “A Crash Course in Good and Bad Controls.”

What is a DAG?



D = education, Y = wages, A = ability

- ▶ Nodes are variables; arrows are direct causal effects; no cycles allowed (“acyclic”).
- ▶ Every **missing arrow is an assumption**: this graph asserts nothing else causes both D and Y .
- ▶ Unobserved variables (like ability) still belong in the graph — being unobservable doesn’t stop something from causing things.

In the graph on the last slide there are two paths from D to Y :

1. $D \rightarrow Y$: the **causal path**. This is the thing we want to measure.
 2. $D \leftarrow A \rightarrow Y$: a **backdoor path**. Association flows along it even though D causes none of it.
-
- ▶ The regression of Y on D alone picks up *both* paths — this is exactly the omitted variables bias formula from Lecture 2, drawn as a picture.
 - ▶ **Identification strategy** = a plan for shutting every backdoor path while leaving the causal path alone.


The grammar: three ways association flows

Every path, however long, is built from three pieces (Mixtape Ch. 3):

Motif	Shape	Default	Conditioning on the middle node
Chain	$A \rightarrow B \rightarrow C$	open	blocks it
Fork	$A \leftarrow B \rightarrow C$	open	blocks it
Collider	$A \rightarrow B \leftarrow C$	blocked	opens it

- ▶ Association flows between two variables if **any** path between them is open. A path is open unless it is blocked at some node along the way.
- ▶ Everything on the coming slides — good controls, bad controls, selection bias — is just these three rows applied to a picture. (The formal name for the full rule is *d-separation*.)
- ▶ The collider row is the one that breaks intuition: it is the only place where controlling for something *creates* association.


Good controls: closing a backdoor



Shading = conditioned on.

- Conditioning on a **confounder** X (regression control, matching, stratification) **blocks** the backdoor path $D \leftarrow X \rightarrow Y$.
- This is the **backdoor criterion** (informally): find a set of observed variables that blocks every backdoor path, condition on them, and the remaining association is causal.
- “Selection on observables” (matching, next month) is exactly the claim that such a set exists *and you have it*.


Colliders: conditioning can open a path



T = talent, L = luck, S = startup gets funded.

- ▶ S is a **collider** on the path $T \rightarrow S \leftarrow L$: two arrows collide into it. A collider **blocks** the path by default.
- ▶ But **conditioning on a collider opens the path**: among *funded* startups, talent and luck are negatively correlated — if a funded founder wasn't talented, they must have been lucky.
- ▶ Talent and luck are independent in the population. The correlation is manufactured by the conditioning. This is **selection bias, drawn as a picture**.


Bad control #1: the post-treatment variable



D = education, M = occupation, Y = wages

- ▶ M is a **mediator**: part of *how* education raises wages is by changing occupations.
- ▶ Controlling for occupation **blocks part of the effect you are trying to measure**. The coefficient on D is no longer the total effect.
- ▶ Rule of thumb with real teeth: **do not control for variables that are themselves outcomes of the treatment**. When in doubt, ask: was this variable determined before or after D ?


Bad control #2: the collider / sample-selection control



D = ad exposure, Y = purchase intent, S = visits
the site

- ▶ Estimating the ad effect **only among site visitors** conditions on a collider: both seeing the ad and wanting the product cause visits.
- ▶ Within visitors, exposure and intent are spuriously related even if the ad does nothing.
- ▶ This is the same disease as Heckman's sample selection (Lecture 11) and the "conditioning on graduation" problem in education studies. **Who is in your sample, and did the treatment help put them there?**

M-bias: even a pre-treatment control can hurt



U_1, U_2 unobserved; X observed **before** treatment.


- The comforting heuristic “pre-treatment controls are always safe” is **false**. Here X is measured before D – and it is a **collider** on the path $D \leftarrow U_1 \rightarrow X \leftarrow U_2 \rightarrow Y$.
- Unconditionally that path is blocked at X ; **controlling for X opens it**, importing bias where there was none.
- Story: D = adopting analytics software, Y = productivity, X = a pre-adoption “tech enthusiasm” survey score driven by managerial ambition (U_1 , which also drives adoption) and workforce skill (U_2 , which also drives productivity).
- Timing is not enough – you have to think about **why the control varies** (Cinelli, Forney & Pearl 9/18

A field guide to controls

Candidate control X	Control for it?	Why
Common cause of D and Y (confounder)	Yes	closes the backdoor
Proxy for an unobserved confounder	Usually yes	closes it partially
Pre-treatment predictor of Y only	Yes (precision)	shrinks SEs, no path to D
Mediator ($D \rightarrow X \rightarrow Y$)	No*	blocks the effect itself
Collider (caused by two paths' worth of causes)	Never	opens a closed path
Descendant of Y	Never	conditioning on the outcome
Instrument (affects D only)	No	amplifies remaining bias

- ▶ Row 3 is the positive surprise: in an A/B test, adjusting for *pre-experiment* outcomes (last month's spend) is pure precision gain — this is what platform experimentation teams do routinely.
- ▶ *Unless the **direct** effect is your estimand and you have an instrument for the mediator — the case study coming next.
- ▶ The full 18-model catalog: Cinelli, Forney & Pearl (2022), "A Crash Course in Good and Bad

Case study: when prices respond to your experiment



T = randomized ad/feature, p = price, q = quantity,
 ξ = demand shock (unobserved)


- ▶ A marketplace A/B test improves product pages. Demand rises — and **treated sellers raise prices**, offsetting part of the sales gain.
- ▶ Often the estimand is the **direct effect**: the demand-curve shift, *holding price constant*.
- ▶ So this time we want to hold the mediator fixed — but p is a **doubly bad control**: post-treatment ($T \rightarrow p$) and correlated with the demand shock ($\xi \rightarrow p$: sellers reprice when demand is hot).

Three regressions, one instrument


With randomized T , positive direct effect, and prices that respond (Jeziorski, Leng & Seiler 2025, working paper):

1. q on T : the **total** effect — demand shift *minus* the price offset. A fine estimand, but it is not the demand shift.
 2. q on T and p : **biased for the direct effect even though T is randomized** — p drags ξ into the regression. Typically the bias shrinks but does not vanish, and it has the same sign as omitting p . The middle answer is still a wrong answer.
 3. The fix: a **price instrument** (a cost shifter that moves p but not demand) recovers the direct effect. Instruments are next week's topic — note that here the instrument is for the **control**, not the treatment.
- Built-in diagnostic: regress p on T . If price responds to your randomized treatment, regressions (1) and (2) answer different questions than you think.
 - Stakes: in supermarket scanner data, not instrumenting price overstates feature-advertising

The research designs of this course, as DAGs



Randomization: coin flip R causes D ; no arrow into D from U .



Instrumental variables: the missing arrows $Z \rightarrow Y$ (exclusion) and $U \rightarrow Z$ (exogeneity) ARE the IV assumptions.

Every design we meet from here on is a claim about missing arrows. When you referee a paper, **draw its DAG** — the contestable assumption is usually an arrow the authors need to be absent.

You don't have to do this in your head

- ▶ **dagitty** (<https://dagitty.net>, also an R package): draw your DAG in the browser and it computes, automatically,
 - every **minimal adjustment set** that identifies your effect (“control for exactly these”), and
 - the **testable implications**: conditional independencies your DAG claims, which you can check in data.A DAG is a falsifiable object, not a vibe.
- ▶ Workflow worth adopting: draw the DAG *before* touching the data; put it in your appendix; let referees argue with the arrows instead of guessing your assumptions.
- ▶ Pearl's machinery goes much further — front-door identification, do-calculus, combining experimental and observational data (“data fusion”) — see Hünermund & Bareinboim (2021) if you want the full toolkit. For this course, backdoors and colliders do 95% of the work.

Where DAGs strain (and economists push back)

- ▶ **Simultaneity**: supply and demand determine price and quantity *jointly* — a cycle, which a DAG cannot draw. Our oldest identification problem (Working 1927, Lecture 6) sits awkwardly in this language.
- ▶ **Heterogeneous effects**: LATE, compliers, and monotonicity (Lecture 7) are natural in potential outcomes and clumsy in DAGs.
- ▶ **Where DAGs shine**: bad controls, sample selection, and transparent communication of assumptions — especially in fields (marketing, IS) where “we controlled for everything” is still treated as a virtue.
- ▶ **Verdict**: use the DAG to decide *what to control for*; use potential outcomes to decide *what your estimate means*.

Coming attraction: when units interfere

- ▶ Everything today assumed my treatment doesn't affect *your* outcome (part of **SUTVA**). Drawn as a DAG: no arrows *between units*.
- ▶ In marketplaces and platforms this fails constantly:
 - Discounting rides for treated riders lengthens wait times for control riders.
 - Promoting treated sellers steals demand from control sellers $\Rightarrow$ the A/B test **overstates** the effect of the promotion.
- ▶ Fixes people use: cluster-randomize markets, switchback designs, exposure modeling. See Wager Chs. 11–12 — and Econometrics II if you want the full treatment.
- ▶ For now: when someone shows you an experiment on a **connected** system, ask what the control group's counterfactual really was.

Recap

- ▶ A DAG is a set of causal assumptions; the missing arrows do the work.
- ▶ Three motifs — chain, fork, collider — are the whole grammar. Backdoor paths are confounding; blocking them is what “controlling for” means.
- ▶ Colliders reverse the logic: conditioning *creates* association. Selection into the sample is a collider, and M-bias shows even pre-treatment controls can be colliders.
- ▶ Bad controls: ask *why* the variable varies, not just *when* it was measured. When in doubt, the field guide — or let dagitty compute the adjustment set.
- ▶ Every design in this course is a claim about missing arrows — draw the picture before you believe the number.

Let's look at the workbook now.

- ▶ Open `inclass_session5.Rmd` (in the L7 - Program Evaluation and Selection folder of the course repo).
- ▶ Reproduce the power calculations, the ten-line randomization inference, and the funded-startup collider.
- ▶ Then work the two problems with a neighbor:
 1. The bad control, live
 2. Design your own experiment

Thanks!
